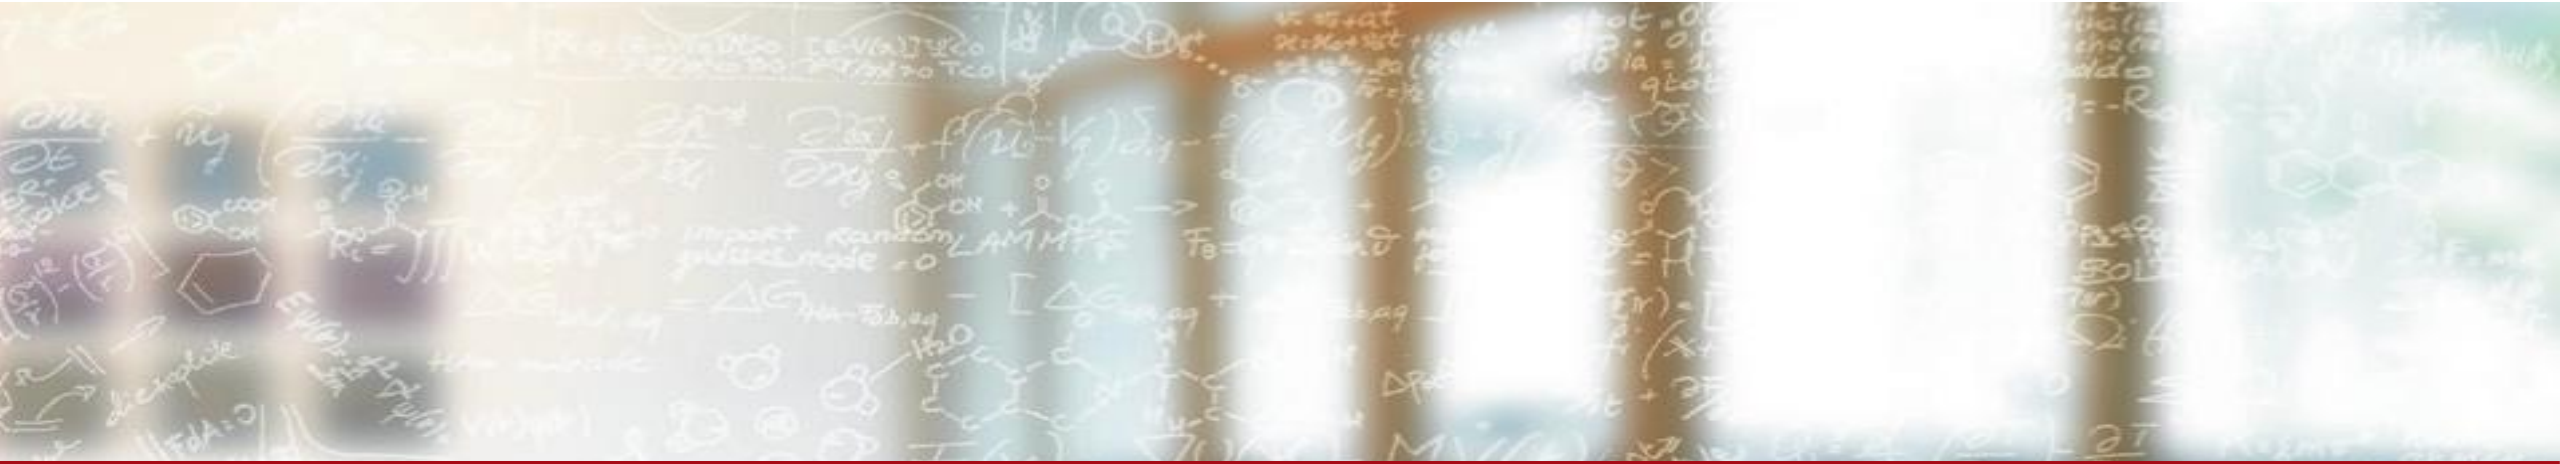




**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



# Onboarding SDSC users at CSCS

## Workshop

Prashanth Kanduri and Lukas Drescher  
CSCS

19th October, 2023

# Structure of the day

## Morning

- Alps overview
- MFA access
- SSH configuration:
  - Daint login nodes and compute nodes
  - Login with VS code
- Running jobs using sbatch
- Conda environment
  - create custom jupyter-kernel
    - shared between Jupyter service, IDE and shell
    - finish by running python script through sbatch

## Afternoon


- Running containers with Sarus on Piz Daint
  - ...using NGC containers for single node and distributed deep learning
  - Large scale training on Piz Daint in MLPerf
- Outlook on Alps
  - New container engine
  - FirecREST for automated workflows
  - High-performance data science with RAPIDS on Clariden

# Containers - key concepts

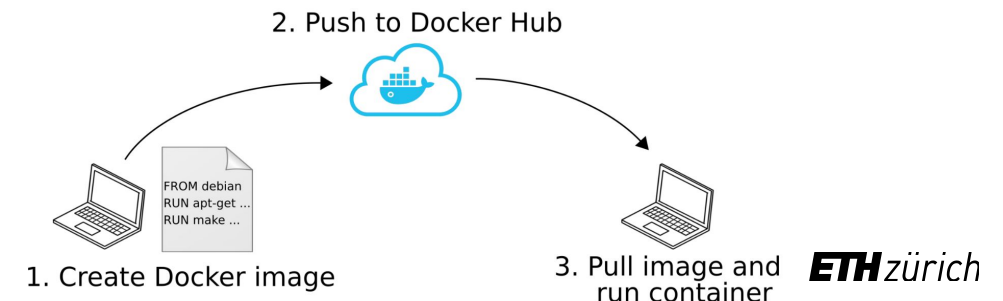
- **Isolated environments** to run applications/services
- **Images** include all software dependencies
  - standalone, executable package including code, runtime libraries, configuration files of an application
  - provides the filesystem and metadata (e.g. environment variables, initial working directory) for a container
- **Container** : a process isolated from the rest of the system through abstractions created by the OS.
  - The level of isolation can be controlled, allowing access to host resources.
  - Its filesystem content comes from an image.
  - “runtime instance” of an image: what the image becomes in memory when executed.
- **Reproducible** (explicit dependencies, consistent environment, immutable infrastructure, host-isol.), **portable, easy to build, quick to test & deploy**

## Linux containers ecosystem

- Linux containers rely on abstraction features (namespaces<sup>1</sup>) provided by the kernel
- Different design decisions and use cases gave rise to several solutions:



<sup>1</sup> “Namespaces in operation, part 1: namespaces overview” at <https://lwn.net/Articles/531114/>



# Using NVIDIA GPUs in Docker

## GPU-accelerated application

- Included/built in the image, along with its runtime dependencies
- NVIDIA provides base images for CUDA, featuring compilers, runtime and accelerated libraries: <https://hub.docker.com/r/nvidia/cuda>
- Quickest way to get a Dockerfile going:  
`FROM nvcr.io/nvidia/cuda`

## GPU driver

- It is tied to the hardware: cannot be part of a portable image!
- Has to be imported upon container creation
- NVIDIA Container Toolkit to the rescue!  
<https://github.com/NVIDIA/nvidia-docker>
- Docker >= 19.03 has native support:  
`docker run --gpus all  
nvcr.io/nvidia/cuda nvidia-smi`

# Docker and HPC: not a good fit

- Security model assumes root privileges
- No integration with workload managers
- Missing support for diskless compute nodes
- Very limited support for kernel bypassing devices (e.g. accelerators and NICs)
- No adequate parallel storage driver



# The Sarus container engine

- Security oriented to HPC systems
- SLURM-aware
- Integrates with HPC infrastructure and software
- OCI-compliant runtime that customizes containers through OCI hooks
  - HW can achieve native performance
- Pulls regular Docker images (e.g. from Docker Hub etc.) and provides Docker-like CLI

→ combines container portability with native HPC performance

# Sarus resources

- [user.cscs.ch](https://user.cscs.ch) → Containers → Sarus → Setup



ETH zürich

- Comprehensive user guide
- Unofficial cheat sheet by Tomas Aliaga
- More materials at [eth-cscs.com/containers-hands-on](https://eth-cscs.com/containers-hands-on) by Alberto Madonna



Basic Commands		
help	List <b>commands</b> and <i>options</i>	
version	Print sarus semantic version	
--debug	Make commands output debug logs	
--verbose	Make commands output info logs	
<div><div></div><div><div></div><div>Sarus is driven by its <i>sarus.json</i> configuration file, typically located at <code>/opt/sarus/&lt;version-number&gt;/etc/sarus.json</code>.</div></div></div>		
Images		
images	sarus images	List local images
pull	sarus pull python:alpine	Pull image from DockerHub
load	sarus load ./debian.tar my_debian	Loads tarball as image
rmi	sarus rmi centos:7	Deletes image
Runtime		
run	sarus run alpine cat /etc/os-release	
--tty	Allocate a pseudo-TTY in the container	
--entrypoint echo hi	Overwrite the default ENTRYPOINT of the image	
--mount=type=bind,src=HD,dst=CD	Mount custom host directories HD into the container at CD	
--mpi	Enable MPI support	
--ssh	Enable SSH in the container	
--centralized-repository	Use centralized repository instead of the local one	
<div><div></div><div><div></div><div>You need to <b>pull</b> the image before you can <b>run</b> the container.</div></div></div>		

on compute node

Useful Links	
📖 documentation	<a href="https://sarus.readthedocs.io">https://sarus.readthedocs.io</a>
</> code	<a href="https://github.com/eth-cscs/sarus/">https://github.com/eth-cscs/sarus/</a>
📦 download	<a href="https://github.com/eth-cscs/sarus/releases">https://github.com/eth-cscs/sarus/releases</a>

on compute node (prev. salloc)

Examples on Localhost	
GPU	sarus run ethcscs/cudasamples:9.2
nbody	/usr/local/cuda/samples/bin/x86_64/linux/release/nbody -benchmark -fp64 -numbodies=2000

on login node, through SLURM

Examples on HPC	
GPU	srun -C gpu -N1 -t1 sarus run
nbody	ethcscs/cudasamples:9.2 /usr/local/cuda/samples/bin/x86_64/linux/release/nbody -benchmark -fp64 -numbodies=200000
MC	srun -C gpu -N2 -t2 sarus run --mpi
all2all	ethcscs/osu-mb:5.3.2-mpich3.1.4 ./osu_latency
❗ Find more examples in the <a href="#">Sarus Cookbook</a> .	

# Practical part - developing with Sarus on Piz Daint

Go to <https://github.com/lukasgd/SDSC-user-onboarding>



# Distributed Training

Instructions for Horovod (alternative to DDP):

- [user.cscs.ch](https://user.cscs.ch) → Data Science → [PyTorch](#)

Large-scale distributed training comes in different flavors

- data-parallel is most accessible
- pipeline-, model-parallel more advanced
- Experience from MLPerf HPC benchmark
  - <https://mlcommons.org/en/training-hpc-07/>, [hpc-10/](#)
  - significant speedups possible, e.g. DeepCAM with 9 TB requiring ~300 GPU hours converges in ~20 mins on 1000 GPUs!
  - more details presented at [MLHPC SC'21](#)

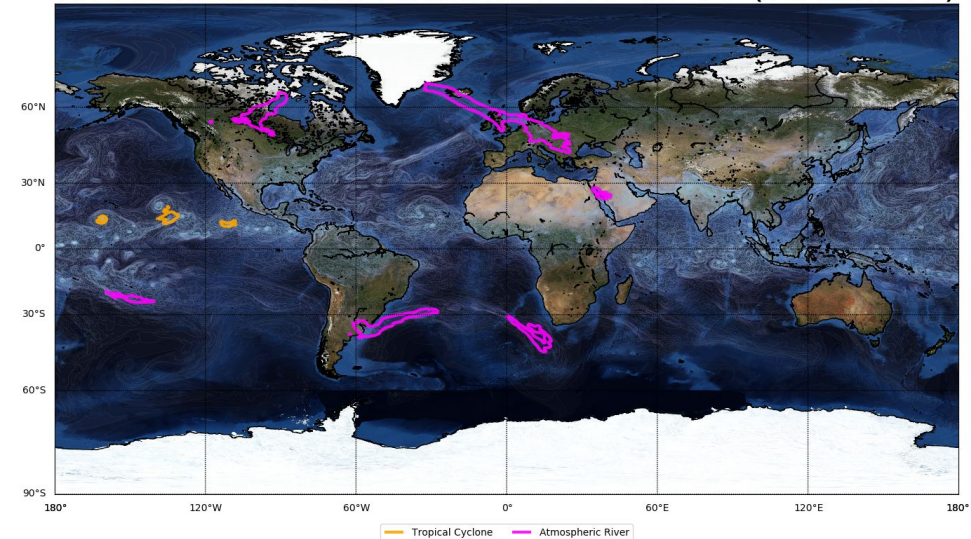
Considerations for scaling SGD with data-parallelism

- Optimize on a single GPU first
  - data loader - check GPU utilization, compare to loading synthetic data, be aware of kernel launch overhead
- Then scale horizontally
  - aim for ideal weak scaling, tune communication
- Be aware of convergence (#epochs) as a function of batch size

## DeepCAM

- Gordon-Bell prize paper at SC18 (Kurth et al)
- convolutional encoder-decoder segmentation model
- identifies extreme weather phenomena: atmospheric rivers and tropical cyclones (3 per-pixel classes)
- 8.8 TB climate simulation data (HDF5), 0.25° spatial, 3h temporal resolution, shape (768, 1152, 16)
- Reference implementation in PyTorch

Extreme Weather Patterns 2007-05-07 (stream 02)





# Outlook on Alps

## Flash-based scratch FS

### New container engine on Alps

- based on enroot/pyxis
- more user-friendly
  - TOML format
  - Infrastructure-as-Code
  - closer integration with SLURM

```
$ cat env.toml
image = "/iopsstor/scratch/.../.enroot/my_images/NGC-pytorch-23.09-py3.sqsh"
mounts = ["/iopsstor/scratch/cscs/lukasd:/scratch", "/users/lukasd"]
workdir = "/iopsstor/scratch/cscs/lukasd/path/to/repo"
writable = true
```

```
$ srun -ul --environment=$(pwd)/env.toml --gpus-per-task=1 bash -c '...'
```

### Clariden

- Current testbed for the ML platform

#### Node configuration

- 128-core CPU
- > 500 GB RAM
- 4 A100 GPUs per node

Enables high-performance data science with RAPIDS:

<https://github.com/lukasgd/SDSC-user-onboarding/tree/main/clariden>



**REST API** for HPC systems  
<https://firecrest-api.cscs.ch/>



**Accessibility:** opens web access to HPC from any device



**Adaptability:** modular design to support diverse HPC ecosystems



**Security:** provides multiple authorization control layers



**Programmability:** uses standard interface for a simplified automation

# Take Home Messages

- It is very easy to set up persistent Conda environments on CSCS machines
  - They are reusable in the shell, IDE and the Jupyterlab service
  - Use miniconda to create custom development environments
  - It is also possible to attach IDEs (like VS Code) to compute nodes
  - Writing multi-gpu ML programs is easy, with DDP/Horovod/etc, then launch with `SBATCH` scripts
- **Sarus**, our high performance container engine allows using Docker images pushed to Dockerhub on compute nodes at CSCS
  - Start with [Nvidia NGC containers](#) for key packages like PyTorch/TensorFlow/etc as a basis
  - Containers are used for **CI on CSCS hardware** for development projects hosted on Git repos
- The **future ALPS platform** will be a **containers-first ecosystem** where replicating environments is designed to be seamless and user-friendly
  - **Full flash storage** for home and scratch for fast I/O
  - **Powerful CPU+GPU** hardware with **large memory** and **high bandwidths**
  - Dedicated cluster and service offering for the ML/AI user community
  - Customized workflows and services can be built with **FirecREST** ([products.cscs.ch](https://products.cscs.ch))


... also, we are just one floor away!



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



**Thank you for your attention.**